Historical citation network: included papers and their references

Focus view: all 45 included papers + the 25 most frequently cited external references | 199 direct links shown

Complete reconstructed network: 1,924 nodes and 2,237 links; cited-reference years 1947–2026

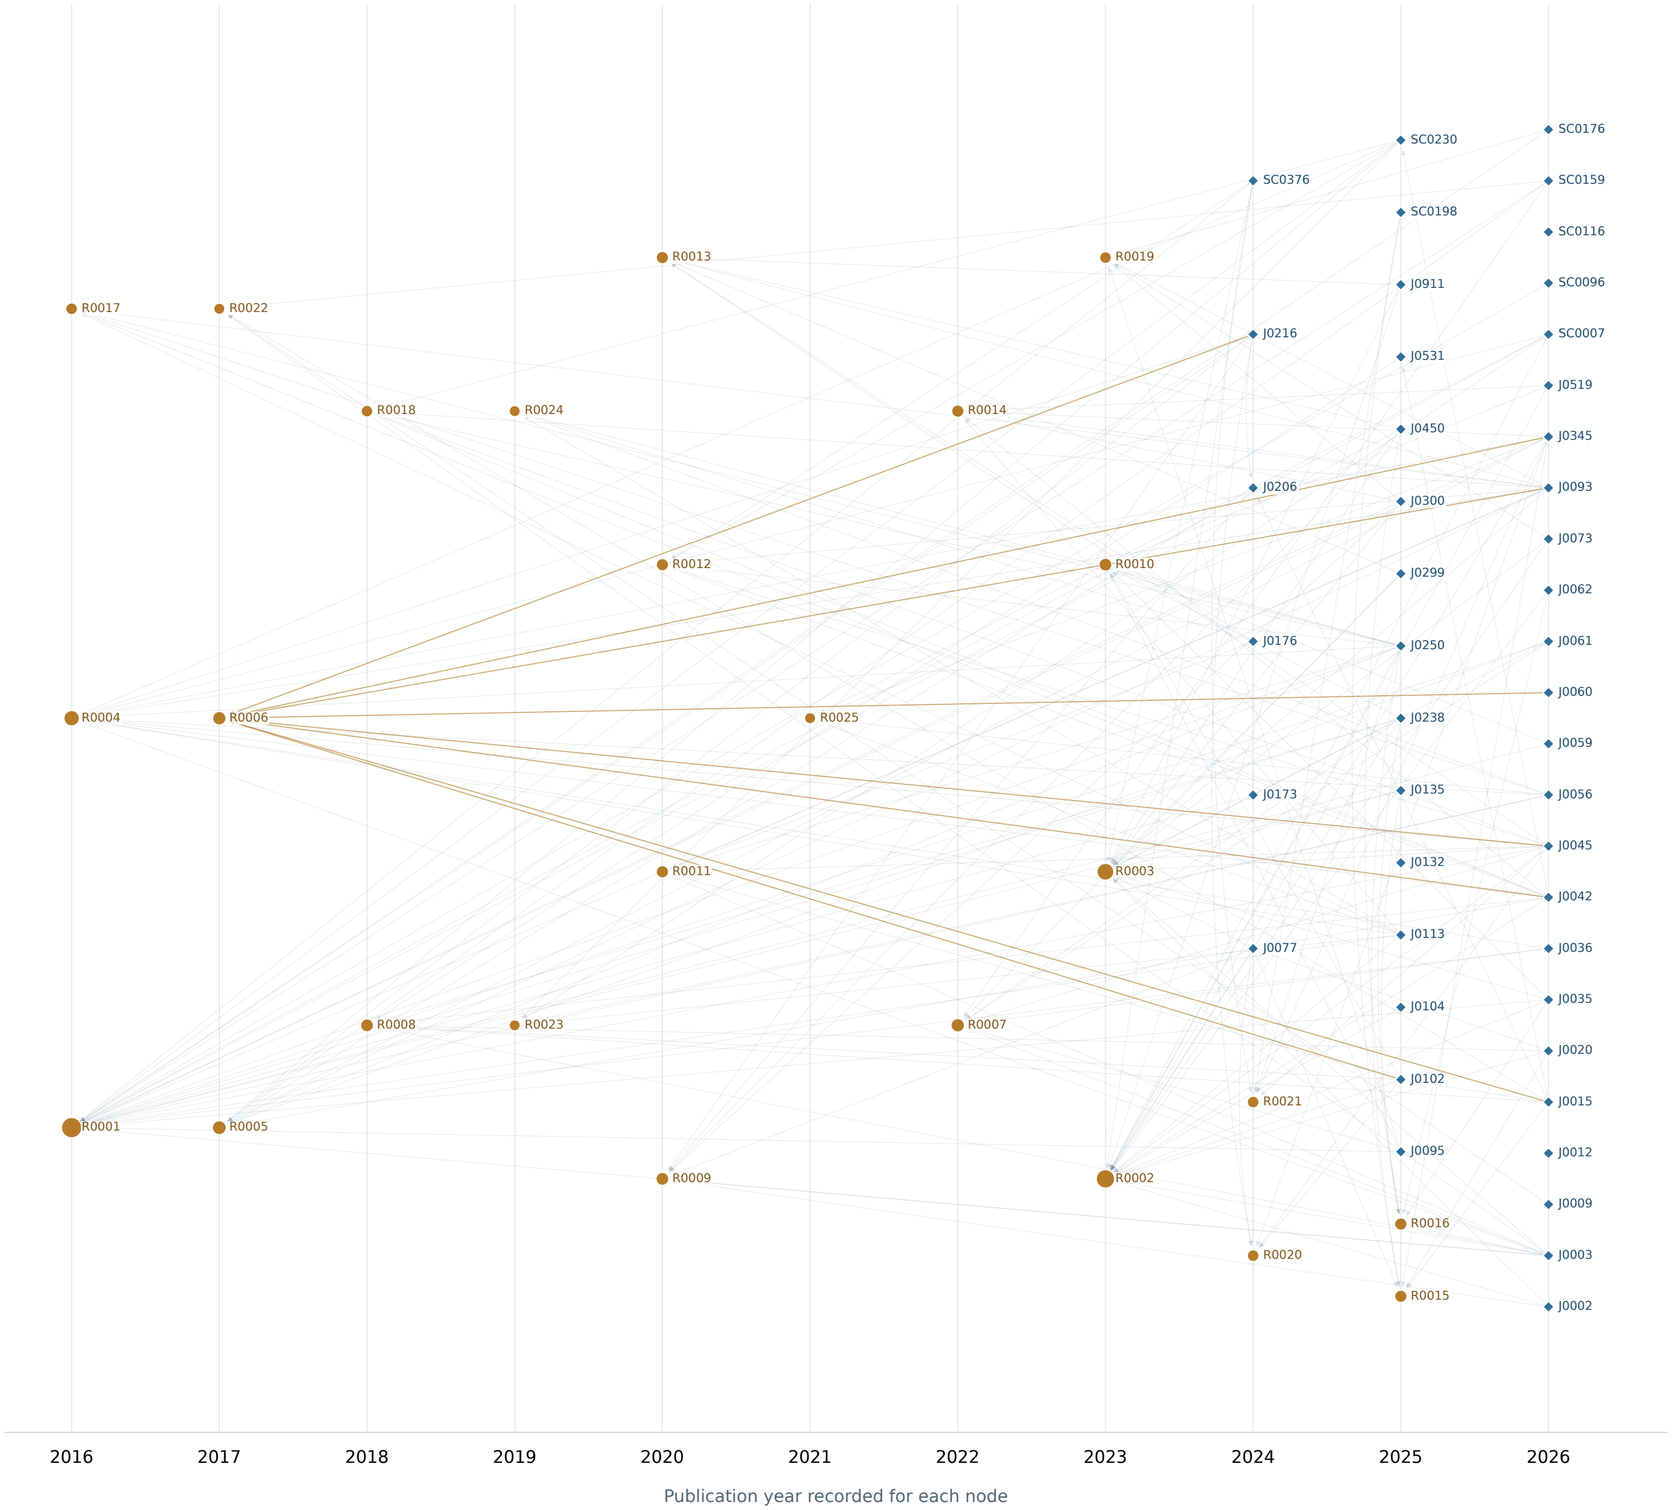

Publication year recorded for each node

Blue diamonds: included papers (review IDs). Gold circles: cited references; circle area increases with the number of included papers citing the work.

Arrows run from citing paper to cited work. Gold links highlight the 2017 reference “Attention is all you need”. Vertical placement is for readability.

The complete network and reference evidence are supplied separately. No citation links among external references were inferred. Five source-paper year conflicts and undated references remain in the audit.

CITED REFERENCE KEY

R0001 2016 · cited by 24/45 papers
MIMIC-III, a freely accessible critical care database

R0002 2023 · cited by 19/45 papers
MIMIC-IV, a freely accessible electronic health record dataset

R0003 2023 · cited by 15/45 papers
Large language models encode clinical knowledge

R0004 2016 · cited by 12/45 papers
RETAIN: an interpretable predictive model for healthcare using reverse time attention mechanism

R0005 2017 · cited by 9/45 papers
GRAM: Graph-based attention model for healthcare representation learning

R0006 2017 · cited by 8/45 papers
Attention is all you need

R0007 2022 · cited by 8/45 papers
Chain-Of-Thought Prompting Elicits Reasoning in Large Language Models

R0008 2018 · cited by 7/45 papers
The eICU Collaborative Research Database, a freely available multi-center database for critical care research

R0009 2020 · cited by 7/45 papers
Retrieval-augmented generation for knowledge-intensive NLP tasks

R0010 2023 · cited by 7/45 papers
Gpt-4 technical report

R0011 2020 · cited by 6/45 papers
BEHRT: Transformer for Electronic Health Records

R0012 2020 · cited by 6/45 papers
HiTANet: Hierarchical Time-Aware Attention Networks for Risk Prediction on Electronic Health Records

R0013 2020 · cited by 6/45 papers
Language models are few-shot learners

R0014 2022 · cited by 6/45 papers
Lora: Low-rank adaptation of large language models

R0015 2025 · cited by 6/45 papers
ColaCare: Enhancing Electronic Health Record Modeling through Large Language Model-Driven Multi-Agent Collaboration

R0016 2025 · cited by 6/45 papers
Toward expert-level medical question answering with large language models

R0017 2016 · cited by 5/45 papers
Doctor AI: predicting clinical events via recurrent neural networks

R0018 2018 · cited by 5/45 papers
Interpretable Representation Learning for Healthcare via Capturing Disease Progression through Time

R0019 2023 · cited by 5/45 papers
Large language models in medicine

R0020 2024 · cited by 5/45 papers
Prompting large language models for zero-shot clinical prediction with structured longitudinal electronic health record data

R0021 2024 · cited by 5/45 papers
RAM-EHR: Retrieval Augmentation Meets Clinical Predictions on Electronic Health Records

R0022 2017 · cited by 4/45 papers
On calibration of modern neural networks

R0023 2019 · cited by 4/45 papers
GAMENet: Graph Augmented MEmory Networks for Recommending Medication Combination

R0024 2019 · cited by 4/45 papers
Pre-training of Graph Augmented Transformers for Medication Recommendation

R0025 2021 · cited by 4/45 papers
SafeDrug: Dual molecular graph encoders for recommending effective and safe drug combinations